

# Dependently Typed Programming with Mutable State

Aaron Stump<sup>1</sup>   Evan Austin<sup>2</sup>

<sup>1</sup>Computer Science  
The University of Iowa

<sup>2</sup>Computer Science  
The University of Kansas

U.S. National Science Foundation CAREER grant.

# What Are Dependent Types?

- Indexed datatypes:

<code>&lt;list A n&gt;</code>	instead of	<code>&lt;list A&gt;</code>
<code>&lt;balanced_tree A d&gt;</code>		<code>&lt;tree A&gt;</code>
<code>&lt;lam_t max_var&gt;</code>		<code>lam_t</code>

- Dependent function types:

```
delete : Fun(x:nat) (l:<list nat (S n)>
    (u:{(in x l) = tt})).
    <list nat n>
append : Fun(A:type) (n1 n2:nat)
    (l1:<list A n1>) (l2:<list A n2>).
    <list A (plus n1 n2)>
```

- Computing a type by recursion:

```
printf : Fun(s:format_string).(printf_t s)

(printf_t "%d"++s) => (int -> (printf_t s))
(printf_t "%x"++s) => (ptr -> (printf_t s))
(printf_t []) => unit
```

# Why Dependent Types Matter <sup>1</sup>

## Incrementality

---

<sup>1</sup>Title of invited talk at POPL 2006 by James McKenna.

# Why Dependent Types Matter <sup>1</sup>

Incrementality      Intensionality

---

<sup>1</sup>Title of invited talk at POPL 2006 by James McKenna.

# Incrementality

- Adding verification usually is a big leap.
  - ▶ new specification language (at least first-order logic); and
  - ▶ new proof language(s), or
  - ▶ unpredictable, tricky tools (“you need an expert”).
- Not a big leap with dependent types.
  - ▶ From `<list A>` to `<list A n>` is easier.
  - ▶ Add verification judiciously, “pay as you go”.
- Goal: enable gradual increase in code quality.
  - ▶ Deep verification is at one limit.
  - ▶ Lightweight verification can improve code a lot.

# Intensionality (Policies versus Properties)

- Properties expressing facts about code.
- Policies restrict how code can be used.
- Stating (proving) a property from a policy may be hard.
- Example policies:
  - ▶ Files may not be accessed after they are closed.
  - ▶ Uninitialized array locations may not be read.
  - ▶ Data computed from user's contact list cannot be auto-emailed. <sup>2</sup>.

---

<sup>2</sup>See [Swamy, Chen, and Chugh 2009]

# GURU at a High-Level

- Pure functional language + logical theory.<sup>3</sup>
  - ▶ Includes indexed datatypes, dependent function types.
  - ▶ Terms : Types.
  - ▶ Proofs : Formulas.
- Inspired by Coq/CIC, but with some improvements:
  - ▶ General recursion for terms.
    - ★ Proofs are still sound.
    - ★ Explicit casts instead of conversion  $\Rightarrow$  type equivalence still decidable.
  - ▶ Annotations dropped for type equivalence.
    - ★ Including types, specificationnal (“ghost”) data, and proofs.
    - ★ Avoids problems with equality of proofs.
    - ★ Like Implicit Calculus of Constructions (ICC).
  - ▶ Resource-tracking analysis **[new!]**

---

<sup>3</sup>See [Stump, Deters, Petcher, Schiller, Simpson 2009]

# Functional Modeling for Imperative Abstractions

- I/O, mutable arrays, cyclic structures, etc.
- Do not fit well into pure FP.
- Approach: functional modeling.
  - ▶ Define a pure functional model (e.g., `<list A n>` for arrays).<sup>4</sup>
  - ▶ Model is faithful, but slow.
  - ▶ Use during reasoning.
  - ▶ Replace with imperative code during compilation.
  - ▶ Use *linear* (aka *unique*) types to keep in synch.

---

<sup>4</sup>Cf. [Swierstra and Altenkirch 2007]



# Example: Word-Indexed Mutable Arrays

- **Type:** `<warray A N L>`.
  - ▶ `A` is type of elements.
  - ▶ `N` is length of array.
  - ▶ `L` is list of initialized locations.
- `(new_array A N) : <warray A N []>`.
- **Writing to index `i`:**
  - ▶ requires proof:  $i < N$ .
  - ▶ functional model: consume old array, produce updated one.
  - ▶ imperative implementation: just do the assignment.
  - ▶ array's type changes: `<warray A N i::L>`.
- **Reading from index `i`:**
  - ▶ does not consume array.
  - ▶ requires proof:  $i \in L$ .

## Example: FIFO Queues

- Mutable singly-linked list, with direct pointer to end.
- **Aliasing!**
- GURU approach: *heaplets* (part of heap).

Type	Functional Model	Imperative Implementation
<code>&lt;heaplet A I&gt;</code>	list of aliased values	nothing
<code>&lt;alias I&gt;</code>	index into heaplet <code>I</code>	reference-counted pointer

- Unverified queue:
  - ▶ Just memory safety
  - ▶ 138 lines total (6 lines proof).
- Verified queue:
  - ▶ Prove that `qin`-node has no next-pointer.
  - ▶ Requires reasoning about aliases.
  - ▶ 310 lines total (178 lines proof).

# Resource-Tracking and Memory Management

- Memory deallocated explicitly.
- Resource-tracking analysis ensures safety.
- Different resource types available.
  - ▶ `unowned`: for reference-counted data.
  - ▶ `unique`: for mutable data structures.
  - ▶ `<owned x>`: for *pinning* references.

```
x: unowned  
y: <owned x>
```

Not allowed to consume `x` until `y` is consumed.

Can safely omit inc/dec for `y`.

- GURU: no garbage collection!
- “Garbage Collection: Java Application Servers Achilles’ Heel”<sup>5</sup>

---

<sup>5</sup>[Xian, Srisa-an, Jiang 08]

# Empirical Comparison

Benchmark 1: In array storing  $[0, 2^{20})$ , do binary search for each element.

Benchmark 2: push all words in “War and Peace” through 2 queues.

Mutable Array Test	
Language	Avg Real Time
HASKELL	1.14 s
HASKELL (No GC)	0.45 s
OCAML	0.60 s
OCAML (No GC)	0.54 s
GURU	0.57 s

Queue Test	
Language	Avg Real Time
HASKELL	1.33 s
HASKELL (No GC)	0.60 s
OCAML	0.61 s
OCAML (No GC)	0.38 s
GURU	0.58 s

Compilers: ghc 6.10.4, ocamlpt 3.11.1, gcc 4.3.3  
Machine: 2.67Ghz Intel Xeon, 8 GB mem, Linux 2.6.18

# Current Projects

- `versat`: verified modern SAT solver.
  - ▶ Complex code, uses mutable state.
  - ▶ Not too large.
  - ▶ Simple spec.: learned clauses derivable by resolution from input clauses.
  - ▶ With Duckki Oe, Derek Bruce.
- GOLFSOCK: verified LFSC proof checker.
  - ▶ LFSC = (Edinburgh) Logical Framework with Side Conditions.
  - ▶ My proposal for a meta-language for SMT proofs.
  - ▶ Fast C++ implementation (45% overhead for QF\_IDL, difficulty 0-3).<sup>6</sup>
  - ▶ With Cesare Tinelli, Clark Barrett, Tianyi Liang, Yeting Ge, Andrew Reynolds.
- Implementation in GURU in progress.

---

<sup>6</sup>See [Oe, Stump, Reynolds 2009]

# Current Projects

- `versat`: verified modern SAT solver.
  - ▶ Complex code, uses mutable state.
  - ▶ Not too large.
  - ▶ Simple spec.: learned clauses derivable by resolution from input clauses.
  - ▶ With Duckki Oe, Derek Bruce.
- GOLFSOCK: verified LFSC proof checker.
  - ▶ LFSC = (Edinburgh) Logical Framework with Side Conditions.
  - ▶ My proposal for a meta-language for SMT proofs.
  - ▶ Fast C++ implementation (45% overhead for QF\_IDL, difficulty 0-3).<sup>6</sup>
  - ▶ With Cesare Tinelli, Clark Barrett, Tianyi Liang, Yeting Ge, Andrew Reynolds.
- Implementation in GURU in progress.
- “Eat your own dog food!”

---

<sup>6</sup>See [Oe, Stump, Reynolds 2009]

# Current Projects

- `versat`: verified modern SAT solver.
  - ▶ Complex code, uses mutable state.
  - ▶ Not too large.
  - ▶ Simple spec.: learned clauses derivable by resolution from input clauses.
  - ▶ With Duckki Oe, Derek Bruce.
- GOLFSOCK: verified LFSC proof checker.
  - ▶ LFSC = (Edinburgh) Logical Framework with Side Conditions.
  - ▶ My proposal for a meta-language for SMT proofs.
  - ▶ Fast C++ implementation (45% overhead for QF\_IDL, difficulty 0-3).<sup>6</sup>
  - ▶ With Cesare Tinelli, Clark Barrett, Tianyi Liang, Yeting Ge, Andrew Reynolds.
- Implementation in GURU in progress.
- “Eat your own dog food!”
- Let's eat what we grow.

---

<sup>6</sup>See [Oe, Stump, Reynolds 2009]

# Future Goals

- More imperative abstractions:
  - ▶ Statically reference-counted heaplets.
  - ▶ Doubly-linked lists, hashmaps, etc.
- More automation:
  - ▶ Currently: `hypjoin t t' by p1 ... pn end`<sup>7</sup>.
  - ▶ Extend to first-order formulas?
  - ▶ Goal: understandable, predictable tactics (“no expert needed”).
- (For you) to learn more:
  - ▶ Version 1.0 is close to release:

[www.guru-lang.org](http://www.guru-lang.org)

- ▶ “Verified Programming in Guru” book.

---

<sup>7</sup>See [Petcher, Stump 2009].