

Verified Programming in Guru

Aaron Stump¹ Morgan Deters² Adam Petcher³
Todd Schiller³ Timothy Simpson³

¹Computational Logic Center
CS, The University of Iowa

²LSI, Universitat Politècnica de Catalunya, Spain

³CSE, Washington University in St. Louis

Funding from NSF CAREER.

A Vexing Continuum

Real code

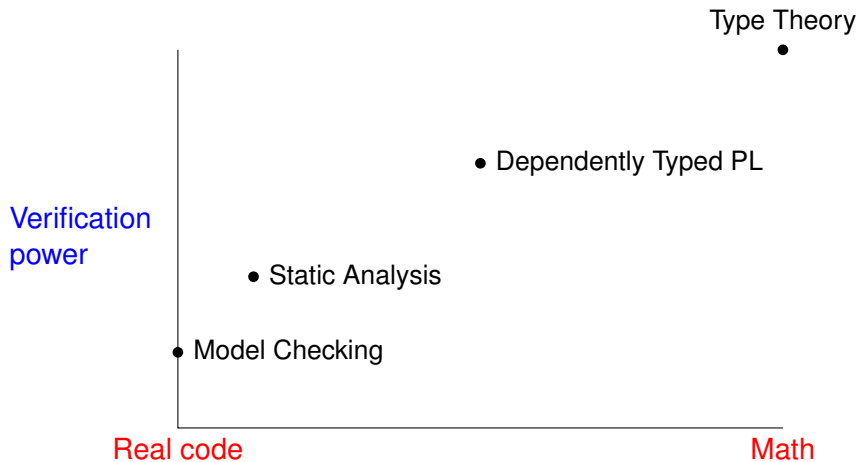
concurrent
imperative
general recursive

Math. functions

sequential
pure
total

Where is your verification method?

Plotting Some Approaches



The GURU Approach

Real code ← GURU Math. functions

General recursion
Dependently typed programs
External theorems about programs
Unaliased mutable state
No concurrency
No aliasing (for mutable state)

Basic GURU Design

- Terms : Types.
- Proofs : Formulas.
- “Full-spectrum” dependency.
 - ▶ Types can contain arbitrary terms.
 - ▶ Definitional equality very weak (no β).
 - ▶ Type checking decidable.
 - ▶ Explicit casts.
- Proofs and types can appear in terms.
 - ▶ computationally irrelevant.
 - ▶ erased by compilation, definitional equality.
- Today:
 - ▶ specificational data.
 - ▶ ownership and memory management.
 - ▶ functional modeling.

Specification Data

- Programmer can designate argument positions `spec`.
 - ▶ for constructors, functions.
 - ▶ can use a `spec x` in a `spec` argument.
 - ▶ also in types, proofs.
 - ▶ nowhere else.
 - ▶ enforced separately from type checking.
- `spec` args erased by compilation [Brady+03], def. equality.
- Improves efficiency, simplifies proofs.

Example: Vector Append

```
Inductive vec : Fun(A:type) (n:nat).type :=  
  vecn : Fun(A:type).<vec A Z>  
| vecc : Fun(A:type) (spec n:nat) (a:A) (l:<vec A n>).  
  <vec A (S n)>.
```

```
vec_append : Fun(A:type) (spec n m:nat)  
  (l1 : <vec A n>) (l2 : <vec A m>).  
  <vec A (plus n m)>
```

Compiled to C: `gvec gvec_append(gtype gA, gvec gl1, gvec gl2);`

```
vec_append_assoc :  
  Forall(A:type) (n1 : nat) (l1 : <vec A n1>  
    (n2 n3 : nat) (l2 : <vec A n2>) (l3 : <vec A n3>).  
  { (vec_append (vec_append l1 l2) l3) =  
    (vec_append l1 (vec_append l2 l3)) }
```

Memory Management in GURU

- Currently, no aliasing.
 - ▶ All data inductive.
 - ▶ Reference graph acyclic.
- Use reference counting, not GC.
- Programs use explicit `inc`, `dec`.
- Static analysis ensures no leaks, no double deletes.
- Analysis runs after type checking.
- Reduce need for `inc/dec` with ownership annotations.

Example: Filling a List

```
fun fill(A:type) (a:A) (n:nat):<list A>.  
  match n with  
    Z => (nil A)  
  | S n' => (cons A a (fill A a n'))  
end.
```

- This type checks, but needs `inc/dec` to compile.
- By default, inputs `unowned` by caller.
- Function must consume each input exactly once.

Compilable Version

```
fun fill(A:type) (a:A) (n:nat):<list A>.  
  match n with  
    Z => dec a (nil A)  
  | S n' => (cons A inc a (fill A a n'))  
end.
```

- `dec a t`: consume reference, evaluate `t`.
- `inc a`: create new reference.
- `n` is consumed by `match`.

A Different Version Using `owned`

```
fun fill(A:type) (owned a:A) (owned n:nat):<list A>.  
  match n with  
    Z => (nil A)  
    | S n' => (cons A inc a (fill A a n'))  
  end.
```

- `a, n` are owned by caller.
- Function must still `inc` for `cons` of `a`.
- No need to `dec a` in `Z` case.
- `match` does not consume owned `n`.
- `n'` automatically owned in second case.

Reference Counting Implementation

- One byte for constructor tag, three for reference count.
- When `refcount = 0`:
 - ▶ put item on per-constructor freelist.
 - ▶ $O(1)$ time.
- When allocating from free list:
 - ▶ `dec subdata`.
 - ▶ $O(d)$ time, where d is arity of constructor.
- Around 4x faster than `malloc/free`.
- For generic code:
 - ▶ pass `int` tags for types.
 - ▶ code for `inc/dec` indexed by tag.

Functional Modeling

- Awkward squad via functional modeling [Swierstra+07].
 - ▶ Identify interface.
 - ▶ Define pure functional model.
 - ▶ Use model for type checking, theorem proving.
 - ▶ Replace during compilation.
 - ▶ Use linear types (`unique`) to ensure equivalence.
- Examples in GURU:
 - ▶ Basic I/O.
 - ▶ 32-bit words with increment.
 - ▶ ASCII characters.
 - ▶ `char`-indexed mutable arrays.

Character-Indexed Mutable Arrays

- **Model** `charvec` **as** `<vec A 128>`.

- **Interface is:**

```
mk_charvec : Fun (A:type) (a:A):unique <charvec A>
```

```
cvget : Fun(A:type) (unique_owned l:<charvec A>)  
      (c:char) : A
```

```
cvupdate : Fun(A:type) (c:char) (a:A)  
          (unique l:<charvec A>) : unique <charvec A>.
```

- `cvget` **does not consume the array.**
- `cvupdate` **does.**

Future Work

- Goal: efficient verified FP with effects.
- So far:
 - ▶ general recursion
 - ▶ mutable structures
 - ▶ good performance via refcounting.
- Next up: aliasing.
 - ▶ idea: maintain a spanning tree of primary pointers.
 - ▶ these have type `unique <aliased A n>`.
 - ▶ `n` is number of outstanding aliases.
 - ▶ to traverse alias, shift primaries/aliases.
 - ▶ use a physical equality to prove equivalent.
 - ▶ eliminate shifting code during compilation.
- Version 1.0 is close to release:

`guru-lang.googlecode.com`